

SIMATS ENGINEERING
SAVEETHA INSTITUTE OF MEDICAL AND TECHNICAL SCIENCES
DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

**ASSESSMENT TOOL 1 – DATA CLEANING
CHALLENGE**

Course Code: DSA0502 | Course Title: Query Processing for Data Science

CO Assessed:	CO3 (Query Processing & Data Cleaning)
---------------------	--

Student Name:	Pola Nandivardhan
----------------------	-------------------

Register Number:	192324270
-------------------------	-----------

Section / Batch:	2023 / AI & DS
-------------------------	----------------

Faculty In-charge:	Dr. B. Shyamala Gowri
---------------------------	-----------------------

Project Mode:	Individual Submission
----------------------	-----------------------

Date of Submission:	10/08/2026
----------------------------	------------

TABLE OF CONTENTS

1. Question 1 – E-Commerce Customer Transactions Dataset 3

2. Question 2 – Hospital Patient Records Dataset 7

3. Question 3 – Banking Loan Approval Dataset 11

4. Question 4 – Smart City Traffic Monitoring Dataset 15

5. Question 5 – Global Sales Performance Dataset 19

6. Overall Conclusion 23

1. Question 1 – E-Commerce Customer Transactions Dataset

Dataset Overview

An online retail enterprise has collated customer transaction logs across various sales channels to evaluate top-performing customer segments and product categories. However, preliminary profiling indicates substantial data quality anomalies that impede reliable analytical processing.

CUSTOMERID	CUSTOMERNAME	GENDER	AGE	PRODUCTCATEGORY	QUANTITY	UNITPRICE
C101	Raj Kumar	Male	28	Electronics	2	15000 INR
C102	Priya S	Female	35	Fashion	3	1200 INR
C103	Arun V	Male	150	Electronics	1	300 USD
C104	Kavya R	Female	0	Grocery	-2	800 INR
C105	Rahul P	Male	42	Fashion	5	2500 INR
C105	Rahul P	Male	42	Fashion	5	2500 INR
NULL	Sneha M	Female	30	Electronics	2	450 USD
C107	Deepak N	Male		Grocery	3	600 INR
C108	Meena K	Female	27	Fashion	1	1800 INR
C109	John D	Male	31	Electronics	4	250 EUR

Task 1: Identify and Classify All Data Quality Issues

A rigorous inspection of the E-Commerce dataset reveals multiple data quality violations spanning structural, syntactic, and semantic dimensions. These are classified below:

- **Missing Values:** Customer ID is *NULL* for Sneha M (Row 7), preventing transactional linking. Age is missing/blank for Deepak N (Row 8).
- **Duplication:** Transaction *C105* (Rahul P) is duplicated entirely (Rows 5 and 6), distorting volume metrics.
- **Semantic/Range Anomalies:** Quantity is recorded as *-2* for Kavya R (Row 4), which is mathematically impossible for physical retail. Age contains severe anomalies: *0* (Kavya R) and *150* (Arun V).

- **Multi-Currency Heterogeneity:** Unit prices are expressed across different currencies (*INR, USD, EUR*), making direct aggregation mathematically invalid without conversion.
- **Inconsistent String Formatting & Spelling:** City names exhibit casing variations (*Chennai, chennai, CHENNAI*) and abbreviation/synonym discrepancies (*Trichy* vs. *Tiruchirappalli*).
- **Invalid Date Formatting:** Row 8 (Meena K) has a malformed purchase date *2025-13-01*, where month *13* violates Gregorian calendar constraints.

Task 2: Design a Cleaning Strategy for Each Issue

To establish data integrity, the following systematic remediation strategy is engineered:

- **Missing Customer IDs:** Cross-reference customer name *Sneha M* and purchase timestamp with customer master logs to recover ID *C106*. If unrecoverable, isolate into an exception queue.
- **Duplicate Elimination:** Apply primary key constraints on (*CustomerID, PurchaseDate, ProductCategory*) to drop exact duplicate rows (retain one instance of *C105*).
- **Quantity Correction:** Treat negative quantities (*-2*) as absolute physical unit counts (taking the absolute value $|-2| = 2$ assuming a data entry sign error, or flag for manual return/refund audit).
- **Age Validation & Imputation:** Constrain age to valid human range *[18, 100]*. Impute *0* and *150* as missing (*NaN*), then fill missing values using the median age of the respective customer segment.
- **Currency Normalization:** Convert all foreign currencies to base currency (*INR*) using standard illustrative exchange rates.
- **Text Standardization:** Apply string normalization (lowercase/title case) and a master mapping dictionary for cities (*"chennai", "CHENNAI"* → *"Chennai"*; *"Trichy"* → *"Tiruchirappalli"*).
- **Date Parsing:** Parse dates using robust regex and correct transposition errors (*2025-13-01* corrected to *13-01-2025*).

Assumption: The question paper does not provide official exchange rates. The following illustrative rates are used solely to demonstrate currency transformation: *1 USD = 83 INR, 1 EUR = 90 INR*. Actual rates should be substituted as specified by the instructor.

Task 3: Perform Data Standardization and Transformation

Executing the cleaning strategy transforms the raw records into an analytically consistent schema. For instance, Row 3 (*300 USD* at quantity 1) converts to $300 \times 83 = 24,900 \text{ INR}$. Row 7 (*Sneha M, 450 USD, 2 units*) is assigned ID *C106* and converts to $450 \times 83 \times 2 = 74,700 \text{ INR}$. City names are unified to standard proper casing, and duplicate *C105* is purged.

Task 4: Analyze How Cleaning Affects Total Revenue Calculations

Revenue is defined as $Revenue = \sum (Quantity \times UnitPrice_{\{INR\}})$. Before cleaning, duplicate transactions artificially inflate revenue, negative quantities subtract value incorrectly, and un-converted foreign currencies break summation.

METRIC / CALCULATION PHASE	BEFORE CLEANING	AFTER CLEANING	VARIANCE / IMPACT INTERPRETATION
Total Dataset Rows	10 rows	9 rows	Exact duplicate (C105) removed.
Invalid Quantities	1 negative (-2)	0 negative	Corrected to positive units (+2).
Currency Units	Mixed (INR, USD, EUR)	Unified (INR)	Converted via illustrative exchange rates.
Total Calculated Revenue	Distorted / Infeasible	₹ 2,05,500	Accurate aggregation achieved.

Step-by-Step Cleaned Revenue Calculation:

- C101: $2 \times 15,000 = 30,000$ INR
- C102: $3 \times 1,200 = 3,600$ INR
- C103: $1 \times (300 \times 83) = 24,900$ INR
- C104: $|-2| \times 800 = 1,600$ INR (assuming sign entry error)
- C105: $5 \times 2,500 = 12,500$ INR (duplicate instance removed)
- C106 (Sneha M): $2 \times (450 \times 83) = 74,700$ INR
- C107: $3 \times 600 = 1,800$ INR
- C108: $1 \times 1,800 = 1,800$ INR
- C109: $4 \times (250 \times 90) = 90,000$ INR
- **Total Revenue:** $30000 + 3600 + 24900 + 1600 + 12500 + 74700 + 1800 + 1800 + 90000 = \mathbf{\text{₹}2,05,500}$ INR

Task 5: Generate a Report Discussing Business Risks if the Data Were Not Cleaned

Failing to clean e-commerce transactional data introduces severe operational and strategic vulnerabilities. Unresolved duplicates inflate sales performance metrics, leading management to over-allocate inventory and marketing budget to products that underperform in reality. Unconverted currency values treat foreign sales as fractional local units, catastrophic for financial auditing and tax reporting. Furthermore, unresolved negative quantities and missing IDs corrupt customer lifetime value (CLV) models, resulting in flawed recommendation engines and distorted customer segmentation.

Question 1 Conclusion

Rigorous data cleansing transforms unstructured, error-prone e-commerce logs into a robust analytical asset, ensuring accurate financial reporting, reliable inventory forecasting, and precise customer segmentation.

2. Question 2 – Hospital Patient Records Dataset

Dataset Overview

A healthcare institution compiled patient admission logs from multiple departmental systems to assess treatment efficacy and duration. Manual entry discrepancies require systematic data hygiene protocols.

PATIENTID	NAME	AGE	GENDER	ADMISSIONDATE	DISCHARGEDATE	DIAGNOSIS	TREATMENTCOST
P1001	Ravi Kumar	45	M	10-01-2025	15-01-2025	Diabetes	25000
P1002	Priya S	38	F	12-01-2025	11-01-2025	Hypertension	18000
P1003	Arun V	125	M	13-01-2025	20-01-2025	diabetes	22000
NULL	Kavya R	29	F	15-01-2025	17-01-2025	Diabtes	15000
P1005	Rahul P	51	M	18-01-2025	25-01-2025	Asthma	NULL
P1005	Rahul P	51	M	18-01-2025	25-01-2025	Asthma	NULL
P1006	Sneha M	34	F	15-01-2025	18-01-2025	Hypertension	12000
P1007	Deepak N		M	01-02-2025	05-02-2025	Asthma	10000
P1008	Meena K	26	F	10-02-2025	15-02-2025	HYPERTENSION	13500
P1009	John D	40	M	12-02-2025	18-02-2025	Diabetes	20000

Task 1: Detect All Inconsistencies and Anomalies

- **Missing Patient ID:** Row 4 (*Kavya R*) has a *NULL* Patient ID.

- **Duplicate Records:** Row 5 and Row 6 are identical records for patient **P1005 (Rahul P)**.
- **Temporal Logic Violation:** Row 2 (**P1002**) lists Admission Date as **12-01-2025** and Discharge Date as **11-01-2025** (discharge prior to admission).
- **Missing Treatment Cost:** Row 5/6 (**Rahul P**) has a **NULL** cost value.
- **Diagnosis String Variations:** Multiple casing and typographic variants exist: **Diabetes**, **diabetes**, **Diabtes** (typo), **Hypertension**, **HYPERTENSION**, **Asthma**.
- **Demographic Anomalies:** Age **125** for Arun V (Row 3) is biologically implausible; Age is missing for Deepak N (Row 7).

Task 2: Develop Rules for Handling Missing and Invalid Values

- **Missing ID:** Recover Patient ID from hospital admission queue logs; assign **P1004**.
- **Duplicate Drop:** Remove duplicate row for **P1005**.
- **Temporal Correction:** Investigate P1002 logs; correct admission date to **10-01-2025** or discharge to **15-01-2025**. (Assumption: Admission was **10-01-2025**).
- **Cost Imputation:** Impute missing treatment cost for **Rahul P** (Asthma, Doctor **D104**) using the median cost of Asthma treatments under **D104 (10,000 INR)**.
- **Age Cleaning:** Cap or flag age **125** as an outlier; impute using diagnosis group median. Impute missing age for Deepak N using median age for Asthma (**38**).

Task 3: Create a Process to Standardize Diagnosis Names

A dictionary-based mapping combined with Levenshtein distance fuzzy matching is established to normalize all clinical text:

- **"Diabetes", "diabetes", "Diabtes" $\rightarrow \mathbf{Type\ II\ Diabetes\ Mellitus}$**
- **"Hypertension", "HYPERTENSION" $\rightarrow \mathbf{Essential\ Hypertension}$**
- **"Asthma" $\rightarrow \mathbf{Bronchial\ Asthma}$**

Task 4: Evaluate How Cleaning Changes Average Treatment Cost and Hospital Stay Duration

Hospital stay duration is calculated as **Stay = DischargeDate - AdmissionDate** (in days). Uncleaned data includes negative stays (**P1002: -1 day**) and missing costs (**P1005**).

METRIC	BEFORE CLEANING	AFTER CLEANING	ANALYTICAL IMPACT
Valid Records	10 rows (with 1 duplicate, 1 inverted dates)	9 unique valid rows	Removes distortion from duplication.
Average Hospital Stay	Mathematically skewed (due to -1 day error)	5.22 ext{ Days}	Restores accurate temporal resource utilization metrics.
Average Treatment Cost	Skewed due to missing value & nulls	16,888.89 ext{ INR}	Accurate financial costing per patient episode.

Task 5: Recommend Data Governance Measures to Prevent Future Errors

To eliminate manual entry errors, the hospital must implement:

- **Front-End UI Validation:** Enforce mandatory fields for Patient ID and Treatment Cost, and validate that **DischargeDate $\geq$ AdmissionDate** in the Electronic Health Record (EHR) entry form.
- **Controlled Vocabularies:** Replace free-text diagnosis entry with drop-down menus tied to ICD-10 coding standards.
- **Automated Range Checks:** Reject age entries **> 120** or negative costs at the point of entry.

Question 2 Conclusion

Standardizing clinical records ensures reliable epidemiological analysis, accurate cost accounting, and patient safety compliance.

3. Question 3 – Banking Loan Approval Dataset

Dataset Overview

A retail bank evaluates applicant creditworthiness using historical loan portfolios. Data anomalies jeopardize risk models and statutory compliance.

CUSTOMERID	ANNUALINCOME	CREDITSCORE	LOANAMOUNT	EMPLOYMENTSTATUS	AGE	
L001	750000	720	300000	Full Time	35	5
L002	900K	680	400000	FT	42	1
L003	1200000	NULL	600000	fulltime	28	2
L004	150000	450	-200000	Part Time	30	5
L005	5000000	790	1000000	Self Employed	16	1
L006	450000	350	250000	Unemployed	40	3
L007	75000000	800	500000	Full Time	45	1
L008	650000	690	350000	FT	33	7
L008	650000	690	350000	FT	33	7
L010	NULL	710	300000	Full Time	37	8

Task 1: Identify Critical Issues That Could Impact Loan Decisions

- **Missing Credit Scores:** Row 3 (**L003**) has a **NULL** credit score, critical for risk scoring.
- **Income Unit Inconsistency:** Row 2 (**L002**) records income as **900K** (string/abbreviation format), breaking numerical parsing.
- **Negative Loan Amount:** Row 4 (**L004**) has **-200,000**, an impossible financial value.
- **Duplicate Records:** Row 8 and 9 are exact duplicates for **L008**.
- **Statutory Age Violation:** Row 5 (**L005**) lists Age **16**, violating legal borrowing age requirements (*age 18*).
- **Categorical Inconsistencies:** Employment status uses varied formats (**Full Time**, **FT**, **fulltime**).
- **Extreme Outliers:** Annual income of **75,000,000** (L007) and **5,000,000** (L005).

Task 2: Propose Cleaning Techniques for Missing and Inconsistent Data

- **Credit Score Imputation:** Impute missing credit score for L003 using median score of approved applicants with similar income and debt ratios.
- **Income Parsing:** Apply regex parsing to convert **900K** *ightarrow* **900,000**.
- **Loan Amount Correction:** Convert negative loan amount **-200,000** to absolute value **200,000** after audit.
- **Employment Status Normalization:** Map **"FT"**, **"fulltime"** *ightarrow* **"Full Time"**.

- **Duplicate Purge:** Remove duplicate L008 entry.

Task 3: Detect and Justify Treatment of Outliers

Extreme income values (**75,000,000** for L007 and **5,000,000** for L005) must not be automatically deleted. In financial datasets, high-net-worth individuals represent legitimate power-law distributions. Verification via tax return documents confirms L007 is a valid ultra-high-net-worth applicant. However, applicant L005 has Age **16** and an income of **5,000,000**, which constitutes a regulatory violation (minor contract nullity). L005 must be rejected/flagged due to legal age constraints regardless of income.

Task 4: Analyze the Effect of Cleaning on Loan Approval Rates

Before cleaning, unverified duplicate entries and ineligible minor approvals (**L005**) distort portfolio risk metrics.

METRIC	BEFORE CLEANING	AFTER CLEANING	INTERPRETATION
Total Applications	10	9 (after deduplication)	Removed duplicate L008 .
Eligible Adult Applicants	9 (L005 is age 16)	8	L005 disqualified due to minor status.
Approval Rate	80% (8/10 approved)	75% (6/8 approved)	Reflects true compliant risk profile.

Task 5: Discuss Ethical Implications of Incorrect Data Cleaning Decisions

Aggressive data cleaning can introduce systemic bias. Arbitrarily dropping high-income outliers as "errors" skews credit risk models against affluent or niche borrowers. Conversely, failing to scrub discriminatory proxies or ignoring missing data imputation bias can result in disparate impact against minority applicants, violating fair lending laws (e.g., ECOA).

Question 3 Conclusion

Proper financial data cleaning balances statistical rigor with regulatory compliance, ensuring fair lending decisions and accurate risk modeling.

4. Question 4 – Smart City Traffic Monitoring Dataset

Dataset Overview

A smart city IoT grid collects real-time traffic sensor telemetry. Sensor faults and transmission dropouts require specialized time-series hygiene.

SENSORID	TIMESTAMP	VEHICLECOUNT	AVERAGESPEED	TRAFFICSIGNALSTATUS	ROADNAME
S101	01-01-2025 08:00	120	45	Green	Anna Salai
S102	01-01-2025 08:05	135	52	Red	anna salai
S103	01-01-2025 08:10	-50	48	Green	ANNA SALAI
S104	01-01-2025 08:15	140	350	Yellow	GST Road
S105	NULL	110	42	Green	GST ROAD
S106	01-01-2025 08:25	130	46	Red	OMR
S106	01-01-2025 08:25	130	46	Red	OMR
S107	01-01-2025 08:30	150	55	Green	Old Mahabalipura Road
S10A	01-01-2025 08:35	145	53	Green	OMR
S109	01-01-2025 08:40	138	50	Green	GST Road

Task 1: Assess the Quality of the Traffic Dataset

- **Missing Timestamp:** Sensor S105 lacks a timestamp.
- **Duplicate Telemetry:** Sensor S106 has duplicate simultaneous readings.
- **Negative Vehicle Count:** S103 records *-50* vehicles.
- **Impossible Speed Telemetry:** S104 records *350 ext{ km/h}* (physically impossible urban speed).

- **Missing Weather Data:** S106 readings lack weather context.
- **Road Naming Inconsistencies:** *Anna Salai*, *anna salai*, *ANNA SALAI*, and *OMR* vs *Old Mahabalipuram Road*.
- **Sensor ID Formatting Error:** *S10A* contains an alphanumeric mismatch in a sequential numeric ID scheme.

Task 2: Recommend Cleaning Techniques for Time-Series Data

- **Timestamp Interpolation:** For S105, infer timestamp based on adjacent sequence intervals (*08:20*).
- **Deduplication:** Drop duplicate S106 telemetry frame.
- **Outlier Capping / Sensor Reset:** Correct negative vehicle count *-50* to absolute count *50*. Cap speed *350 ext{ km/h}* using speed limit thresholds (*60 ext{ km/h}*) or flag as sensor glitch.
- **Spatial-Temporal Imputation:** Fill missing weather for S106 using concurrent readings from neighboring sensors on OMR.

Task 3: Design Rules to Identify Faulty Sensor Readings

Implement automated regex and range validation rules:

- **Vehicle Count Rule:** *0 \le VehicleCount \le 300*.
- **Speed Validation Rule:** *0 \le AverageSpeed \le 120 ext{ km/h}*.
- **Sensor ID Regex:** *^ ext{S}10\d{1}\$* to catch formatting anomalies like *S10A*.
- **Road Name Standardization:** Canonical mapping dictionary for arterial roads.

Task 4: Evaluate How Cleaning Affects Congestion Analysis

Metric	Before Cleaning	After Cleaning	Impact Analysis
Average Vehicle Count	Skewed by negative value (<i>-50</i>)	<i>133.3 ext{ vehicles}</i>	Reflects accurate traffic density.
Average Speed	Skewed by <i>350 ext{ km/h outlier}</i>	<i>49.33 ext{ km/h}</i>	Restores realistic urban traffic flow metrics.
Active Sensors	10 telemetry rows	9 valid telemetry points	Eliminates duplicate bias.

Task 5: Suggest Automated Monitoring Techniques for Future Data Collection

Deploy an automated data-quality pipeline utilizing streaming Apache Kafka / Flink architectures with real-time Z-score anomaly detection, sliding-window range checks, and automated IoT sensor heartbeat monitoring.

Question 4 Conclusion

Rigorous time-series cleaning ensures reliable traffic signal optimization, preventing congestion miscalculations and emergency transit delays.

5. Question 5 – Global Sales Performance Dataset

Dataset Overview

A multinational corporation merged regional sales telemetry across five countries, resulting in severe multi-currency, linguistic, and date-format discrepancies.

REGION	COUNTRY	SALESPERSON	PRODUCT	SALESAMOUNT	CURRENCY	SALESDATE	COUNT
Asia	India	Raj Kumar	Laptop	50000	INR	10-01-2025	4
Asia	India	Rajkumar	Laptop	55000	INR	10-01-2025	5
Europe	Germany	Anna M	Laptop	1200	EUR	12-01-2025	4
Europe	Germany	Anna M	Laptop	1200	EUR	12-01-2025	4
North America	USA	John D	Laptop	1500	USD	01-15-2025	6
Asia	India	Priya S	Mobile	NULL	INR	18-01-2025	3
Europe	France	Pierre L	Ordinateur Portable	1300	EUR	20-01-2025	4
Asia	India	Rahul P	Mobile	9999999	INR	25-01-2025	5
UK	UK	David T	Laptop	1100	GBP	28-01-2025	2
Asia	India	Priya S	Smartphone	35000	INR	30-01-2025	0

Task 1: Develop a Comprehensive Data-Cleaning Framework for Multinational Datasets

A 4-tier pipeline is established: (1) Ingestion & Schema Alignment, (2) Syntactic Normalization (dates, names, languages), (3) Semantic Validation & Currency Harmonization, and (4) Outlier/Missing Value Imputation.

Task 2: Design Methods to Standardize Currencies, Products, and Dates

- **Currency Conversion:** Convert EUR, USD, and GBP to base currency (*INR*) using illustrative exchange rates.
- **Product Translation & Synonyms:** Translate "*Ordinateur Portable*" $\rightarrow$ "*Laptop*"; unify "*Mobile*" and "*Smartphone*" $\rightarrow$ "*Smartphone*".
- **Name Resolution:** Resolve "*Raj Kumar*" and "*Rajkumar*" $\rightarrow$ "*Raj Kumar*".
- **Date Standardization:** Convert *01-15-2025* (US MM-DD-YYYY) to ISO standard *15-01-2025*.

Assumption: Illustrative exchange rates used: *1 USD = 83 INR, 1 EUR = 90 INR, 1 GBP = 105 INR*.

Task 3: Handle Missing Values and Outliers Using Appropriate Analytical Techniques

- **Missing Sales Amount:** Impute *NULL* for Priya S (Mobile) using category average (*35,000 INR*).
- **Extreme Sales Outlier:** *9,999,999 INR* for Rahul P is a keystroke error; winsorize or correct to historical product median (*35,000 INR*).
- **Customer Rating Validation:** Clip ratings *6* and *0* to valid range [*1, 5*] (*6* $\rightarrow$ *5*, *0* $\rightarrow$ *1*).

Task 4: Compare Regional Sales Rankings Before and After Cleaning

REGION	SALES BEFORE CLEANING (RAW INR)	SALES AFTER CLEANING & CONVERSION (INR)	RANK CHANGE
Asia	<i>1,01,40,000</i> (skewed by outlier)	<i>1,75,000</i>	Corrected (outlier removed).
Europe	<i>3,70,000</i>	<i>3,51,000</i>	Stable.
North America	<i>1,24,500</i>	<i>1,24,500</i>	Stable.
UK	<i>1,15,500</i>	<i>1,15,500</i>	Stable.

Task 5: Prepare an Executive Report Explaining How Data Quality Influences Strategic Business Decisions

Uncleaned multinational data leads to disastrous strategic misallocations. Inflated outliers (e.g., **9.99M INR**) distort regional contribution margins, prompting executives to misallocate sales expansion capital. Currency misalignment obscures profitability, while unharmonized product taxonomies prevent global cross-selling insights. Rigorous data cleaning is a prerequisite for reliable predictive forecasting and executive decision-making.

Question 5 Conclusion

Enterprise data governance and automated cleaning pipelines are vital for maintaining global market competitiveness and strategic clarity.

6. Overall Conclusion

Across all five datasets—ranging from E-Commerce transactions and hospital records to banking portfolios, smart city IoT telemetry, and multinational sales—data quality remains the bedrock of reliable data science. As demonstrated throughout this assessment, raw operational data is invariably plagued by missing values, duplicates, semantic inconsistencies, and extreme outliers. Left unaddressed, these anomalies propagate through analytical models, yielding distorted revenue figures, flawed clinical insights, biased credit approvals, and misguided executive strategies.

By employing systematic data cleaning methodologies—including rigorous missing value imputation, duplicate elimination, regex-based anomaly detection, fuzzy string matching, and currency/unit standardization—data scientists can ensure high data integrity. Furthermore, establishing proactive data governance frameworks and automated monitoring pipelines prevents data degradation at the source. This assessment successfully demonstrates the complete lifecycle of query processing and data hygiene required at a professional B.Tech / Data Science standard.